
EiKI Summary

Moritz Gerhardt



Table of Contents

1	What is AI?	2
1.1	What is Intelligence?	2
1.2	Characteristics of AI	3
1.3	Foundations of AI	3
2	AI Systems	5
2.1	What is an AI System?	5
2.2	Creating Intelligent Agents	9
3	Uninformed and Informed Search	11
3.1	Problem Solving	11
3.2	Uninformed Search	13
3.3	Heuristics h	16
3.4	Informed Search	17

1 What is AI?

In general, there is no set definition of what artificial intelligence is. The most common definition is:

A system of machines that can perform tasks that normally require human intelligence.

This is flawed in that intelligence itself also does not have a set definition, and neither does human intelligence.

1.1 What is Intelligence?

1.1.1 Touring Test

Intelligence is often defined by the Touring Test. Hereby it is assumed that an entity is intelligent if it cannot be distinguished from another intelligent entity by observing its behavior.

The Touring Test is set up like this:

1. Human interrogator interacts with two entities, A and B. Hereby one of the two is already assumed to be intelligent (another human).
2. If the interrogator cannot distinguish which entity is the assumed intelligent entity, it is assumed that the other entity is intelligent as well.

This test is also flawed as it does not distinguish between different intelligence levels (knowledge, reasoning, language understanding, learning etc.). It also is very subjective to the interrogator. It is not based on an objective metric and therefore the outcome can differ from person to person.

It also assumes that humans are inherently intelligent, which makes it a circular argument.

1.1.2 Chinese Room Argument

The Chinese Room Argument tries to answer the question whether intelligence is the same as intelligent behavior. It assumes that even if a machine behaves intelligently, that does not mean it is actually intelligent.

The argument goes as follows:

1. A person who doesn't know chinese is put into a room. Outside the room there is a person, who can only interact with this person by slipping them notes written in chinese.
2. The person inside the room has detailed instructions as to how to answer the notes, without any translation or understanding.
3. To the person outside of the room it looks like the person inside is able to understand and answer to these notes, therefore the person outside of the room assumes that the person inside knows chinese.

So the person outside of the room assumes intelligence due to behavior, that does not necessitates the intelligence if the answers are already known.

Is a self-driving car intelligent? Is ChatGPT intelligent?

1.1.3 Does it even matter?

The question of what intelligence is is a long, complex and difficult question. However, as it is a more philosophical question it doesn't actually really affect the scientific field of AI. It's a definition, not a basis.

1.2 Characteristics of AI

AI's are often divided into two categories: general AI and narrow AI.

General & Narrow AI

- General (strong) AI is defined so that it can handle any intellectual task
- Narrow (weak) AI is defined so that it has a specific task or domain it works in.

This mirrors the distinction between intelligence and acting intelligent. Currently, most AI's are considered narrow, general AI's are the goal in most research.

An AI should possess the following characteristics:

Adaptability

The ability to improve performance by learning from experience

Autonomous

The ability to perform tasks without constant guidance from a user / expert

Usually an AI is modelled after one of the two following concepts:

Law of Thoughts

The AI should be able to reason about its own actions, arguments and thought processes.
More akin to human intelligence.

Rational Behavior

The AI should be able to determine what the best action is for a given situation to maximize the achievement.
Based on a mathematical model of rationality that achieves the most desirable outcome given the information available.

Rationality has two advantages

- **More General:** For many situations a provable correct option does not exist. The most likely outcome is a better solution.
- **More amenable:** Rationality can be defined, refined and optimized.

However, rationality rarely displays a good model of reality.

Systems that think and behave like humans are often talked about in the scientific field of **Cognitive Science**.

1.3 Foundations of AI

The field of AI is made up by many other fields:

- **Philosophy:** Logic, reasoning, mind as a physical system, foundations of learning, language, rationality
- **Mathematics:** Formal representation and proof algorithms, computation, decidability, tractability, probability
- **Psychology:** Adaptation, phenomena of perception and motor control
- **Economics:** Formal theory of rational decisions, game theory
- **Linguistics:** Knowledge representation, grammar
- **Neuroscience:** Physical substrate for mental processes

- **Control theory:** homeostatic systems, stability, optimal agent design

2 AI Systems

2.1 What is an AI System?

An AI System can be defined as the study of **rational agents** and **their environments**



2.1.1 Environment

- "The surroundings or conditions in which a person, animal, or plant lives or operates" - Oxford Dictionary
- In AI: The surrounding of an (AI) agent, where the agent operates
- Does not have to be real - Can be artificial
- Example:
 - Selfdriving cars: Street, traffic, weather, road signs, ...
 - Chess: Board, pieces,...

Characteristics of Environments

Discrete vs. Continous

Discrete: Environment has countable number of distinct, well defined states

Continous: Not discrete - Uncountable number of states

For Example:

- Discrete: Chess - Every state of the board is mathematically determinable and defined
- Continous: Selfdriving car - Practically infinite positions and conditions

Observable vs. Partially Observable / Unobservable

Observable: State is completely determinable at each point in time

Partially Observable: State is only partially determinable or only determinable at specific points in time

Unobservable: State is never determinable

Static vs. Dynamic

Static: Environment does **not** change while the agent is acting

Dynamic: Environment can change while the agent is acting

For Example:

- Static: Jigsaw puzzle - State does not change without the agents action
- Dynamic: Driving - State still changes even when the agent stops acting

Single Agent vs. Multiple Agents

Single Agent: Environment contains only one agents

Multiple Agents: Environment can contain multiple agents

Accessible vs. Inaccessible

Accessible: Agent can obtain complete and accurate information about the state

Inaccessible: Agent cannot obtain complete or only inaccurate information

Deterministic vs. Stochastic / Probabilistic

Deterministic: Next state is completely determined by the current state and the actions of the agents

Probabilistic: Next state is also influenced by other factors

Episodic vs. Non-episodic / Sequential

Episodic: Each **episode** consists of the agent perceiving and then acting - every action is dependent only on the episode


Non-episodic: Actions are also dependent on past memory

These characteristics are important, as the environment specifies the specific needs of the agent:

- Different environments require different agent designs
- Not every algorithm works for a specific environment

2.1.2 Agents

- **Sense:** Perceives its environment
- **Think:** Makes decisions autonomously
- **Act:** Acts upon the environments



Rules of AI Agents

1. Must be able to perceive its environment
2. Observations must be used to make decisions
3. Decisions should be used to act
4. (Action should be rational)

Rational Agents

Maximizes the performance and yield the best positive outcome.

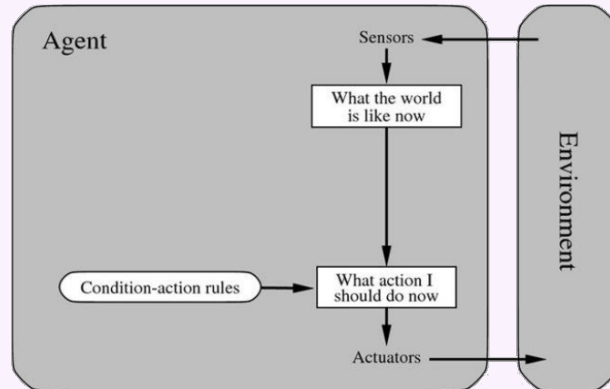
Performance is hereby often measured by a function that evaluates a sequence of actions. This function is task-dependent and cannot be generalized.

2.1.3 Types of Agents

Reflex Agent

Act only on the basis of the current percept, ignores past memory

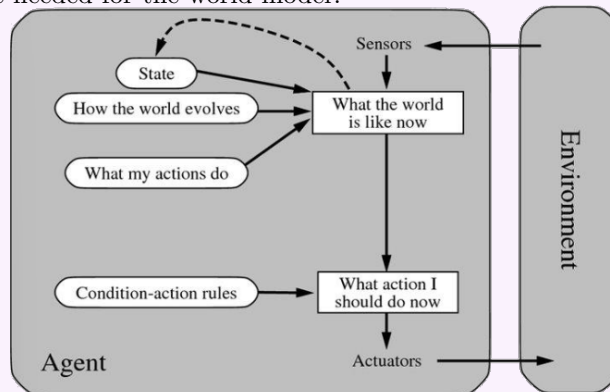
- Implemented through condition-action rules - map state to action
- Problem:
 - Limited decision making
 - No knowledge about anything that cannot be currently perceived
 - Hard to handle in complex environments



Model-Based Agent

Similar decision making to reflex agent, but keep track of the world state


- Input is interpreted and mapped to an internal state representation of the world
- Problem:
 - How do these action affect the internal representation of the world?
 - What details are needed for the world model?



Goal-Based Agent

Essentially a Model-Based Agent with additional functionality that stores desirable states


- Knows what states are desirable and acts towards them
- Problem:
 - Difficult to choose actions if a lot of actions are required to achieve a goal



Utility-Based Agent

Similar to Goal-Based Agents, but instead of providing goals it provides a utility function for rating actions and scenarios based on the desirability of the result

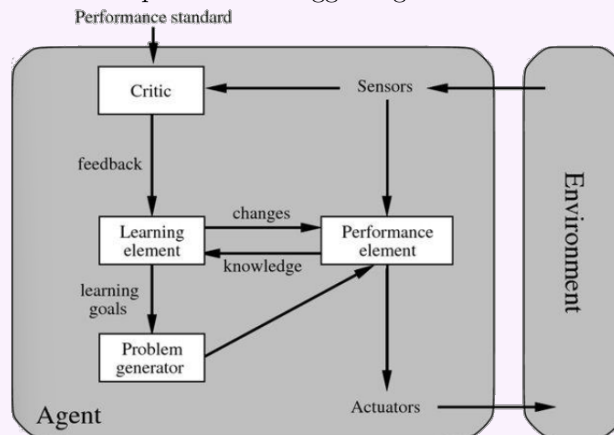
- Goals provide binary distinction, while utility functions provide a continuous measure of desirability
- Can handle selection between two conflicting goals - "Speed or safety"



Learning Agent

Employs additional learning element to gradually improve and become more knowledgeable over time about an environment

- Can learn from past experiences
- Is more robust towards unknown environments
- A learning agent has four conceptual components:
 1. **Learning Element:** Makes improvements by learning from the environment
 2. **Critic:** Gives feedback on the performance of the agent according to a fixed metric
 3. **Performance Element:** Selects the best action according to the critic
 4. **Problem Generator:** Responsible for suggesting actions that will lead to new experiences




2.2 Creating Intelligent Agents

2.2.1 Search Algorithms

Define "finding a good action" as a search problem and use search algorithm to solve it. Most of these search algorithms are tree based, although Bread-First is used often.


2.2.2 Reinforcement Learning

Developed in psychology, Reinforcement Learning essentially is a process of trial and error - Try something out, have a reaction to it, learn whether it was good or bad. Reactions or actions are based on our observations or experiences.



2.2.3 Genetic Algorithms

Inspired by Darwins "Theory of natural selection", this model build multiple different agents and evaluates every one of them. The agent with the highest performance is selected and new models are built based on it. This is done until the performance is good enough.



3 Uninformed and Informed Search


3.1 Problem Solving

3.1.1 Key Terminology

State Space / States

A state is a possible situation in our environment - The State Space is a set of all possible states, reachable from the initial state.

Often visualized as a **state-space graph** that displays all reachable states and its transitions.



State-Space Graph of a robot vacuum cleaning

Transition / Action

A Transition describes possible actions to take to get from one state to another. We only count direct transitions between two states (single actions)

Costs

Transitions often differ in different qualities. We add a "cost" to each action, so we can rate an algorithm on how cost effective it is.

Path

A sequence of states connected by a sequence of actions

Solution

A path that leads from the initial state to a goal state

Optimal Solution

The Solution with the minimal path cost

Problem Components:

1. **State Space and Initial State:**

All possible states and the initial environment as a state

2. Descriptions of Actions:

Function that maps a state to a set of possible actions in this state

3. Goal Test:

Typically a function to test if the current state fulfills the goal

4. Costs:

A cost function that maps actions to costs

An easy way it to add costs of all actions taken

Now that we have defined problems, lets go on to solutions. Most problems can be defined as **planning problems**, in which we start from an initial state and transform it into a desired goal considering future actions and outcomes.

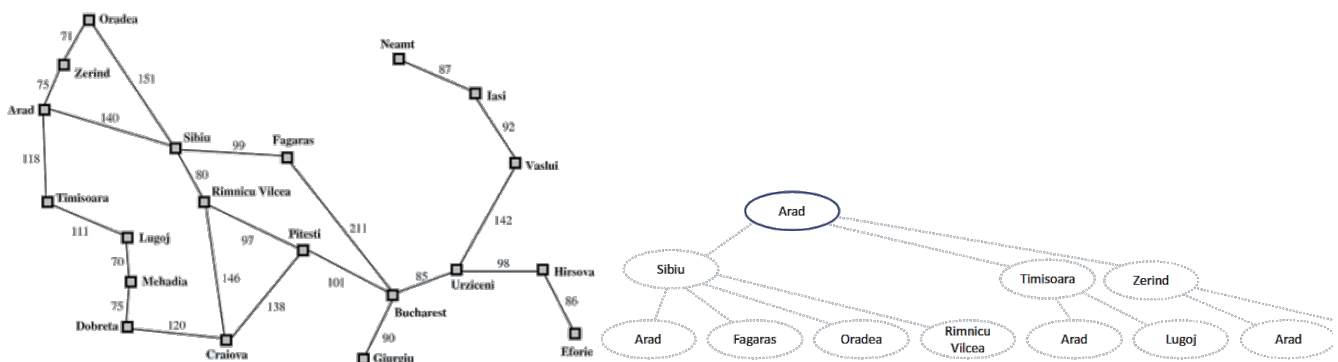
To solve these problems we usually utilize **search** algorithms to find a (optimal) solution in form of a sequence of actions.

3.1.2 Tree Search Algorithms

Hereby we treat the state-space graph as a tree, starting with the initial state. Using this approach we can define iterative or recursive algorithms to search for suitable paths.

Example Tree Search

```
1 Function tree_search(problem, strategy):  
2   initialize search tree with initial state of problem;  
3   While true do  
4     If Node contains goal state then  
5       return solution  
6     Else If No suitable candidates for expansion then  
7       return failure  
8     Else  
9       Expand node and add resulting nodes to the tree
```



Each Node in the search tree is an entire path in the state-space graph. We construct both on demand - only as much as we need (won't get all solutions)

Node

Describes a part of a tree. Includes **state, parent node, taken action, path cost and depth of the tree**

Fringe

Describes the set of nodes at the end of all visited paths

Depth

Number of levels in the search tree

3.2 Uninformed Search

Definition

The uninformed search strategy only has the problem definition, no additional information.

Some algorithms that utilize uninformed search are:

- Uniform-Cost Search (UCS)
- Breadth-First Search (BFS)
- Depth-First Search (DFS)
- Depth-Limited Search (DLS)
- ...

3.2.1 Uniform-Cost Search (UCS)

IN UCS each node is associated with a fixed cost, which accumulate over the path. UCS uses the lowest cumulative cost to find a path.

Only works if each step has a positive cost, as otherwise infinite loops would occur

Complexity: $O(b^{1+\lceil \text{OptimalCost}/\text{eps} \rceil})$ with b being the branching factor (average number of children per node) and eps being the minimal cost for a step

3.2.2 Breadth-First Search

BFS is a special case of UCS, when all costs are equal. It starts at the root and goes through each node of a level before progressing to the next level. BFS stops as soon as it finds a solution, while UCS searches for the 'best' solution by lower cost.

Complexity: $O(b^d)$ with b being the branching factor and d being the depth of the search tree


3.2.3 Depth-First Search (DFS)

Depth-First Search starts at the root and continues on one branch until it finds a solution or failure. In case of failure it backtracks to the current parent node and continues on the next branch.


Unlike UCS and BFS DFS is not complete; It fails if the search space is of infinite depth or has loops.

It's also not optimal; More costly solutions may be found before less costly ones.

Complexity: $O(b^m)$ with b being the branching factor and m being the maximum depth of the search tree



BFS: A,B,C,D,E,F,G,H,I



DFS: A,B,D,E,I,F,C,G,H

Criteria	BFS	DFS
Concept	Traversing tree level by level	Traversing tree sub-tree by sub-tree
Data Structure (Queue)	First In First Out (FIFO)	Last In First Out (LIFO)
Time Complexity	$O(\text{Vertices} + \text{Edges})$	$O(\text{Vertices} + \text{Edges})$
Backtracking	No	Yes
Memory	Requires more memory	Less nodes are stored normally (less memory)
Optimality	Yes	Not without modification
Speed	In most cases slower compared to DFS	In most cases faster compared to BFS
When to use	If the target is relatively close to the root node	If the goal state is relatively deep in the tree

Comparison between BFS and DFS algorithms.

3.2.4 Depth-Limited Search (DLS)

DLS is a variant of DFS. Hereby the search is limited to a depth of d . This means that no infinite search can occur, the trade-off being that it might not find all solutions.

Example Depth-Limited Search

```
1 Function depth_limited_search(problem, limit):
2   return recursive_dls(make_node(initial_state(problem)), problem, limit)
3 Function recursive_dls(node, problem, limit):
4   cutoff_occured = false;
5   If goal_test(problem, state(node)) then
6     return node;
7   Else If depth(node) >= limit then
8     return cutoff;
9   Else
10    ForEach successor in expand(node, problem) do
11      result = recursive_dls(successor, problem, limit);
12      If result == cutoff then
13        cutoff_occured = true;
14      Else If result != failure then
15        return result;
16  If cutoff_occured then
17    return cutoff;
18  Else
19    return failure
```

This, of course, is neither complete nor optimal.

Complexity: $O(b^l)$ with b being the branching factor and l being the maximum depth of the search tree

3.2.5 Iterative Deepening Search (IDS)

IDS is a variant of DLS. It works like DLS, but instead of a fixed depth, it iteratively increases the depth until a solution is found.

Example Iterative Deepening Search

```
1 Function iterative_deepening_search(problem):
2   For depth = 0 to  $\infty$  do
3     result = depth_limited_search(problem, depth);
4     If result != cutoff then
5       return result
```

This change makes it complete and optimal.

Complexity: $O(b^m)$ with b being the branching factor and d being the depth of the search tree. This makes its time complexity equal to BFS. However, the space complexity of IDS is $O(bd)$, which is much better than BFS' $O(b^d)$.

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening
Complete?	Yes*	Yes*	No	Yes, if $l \geq d$	Yes
Time	b^d	$b^{1+\lceil OptCost/eps \rceil}$	b^m	b^l	b^d
Space	b^d	$b^{1+\lceil OptCost/eps \rceil}$	bm	bl	bd
Optimal?	Yes*	Yes	No	No	Yes*

Comparison of search strategies.

3.3 Heuristics h

Denotes a "rule of thumb", a rule that may be helpful in solving a problem.

In Tree-Search, a heuristic is a function that estimates the remaining cost from the current node to the goal.

Can go wrong!

Admissible Heuristic

A heuristic is admissible if it never overestimates the actual cost to reach the goal.

$$h(n) \leq h^*(n)$$

if $h^*(n)$ is the true cost to reach the goal from node n

Consistent Heuristic

A heuristic is consistent if for every node n and successor n' generated by any action a :

$$h(n) \leq c(n, a, n') + h(n')$$

if $c(n, a, n')$ is the cost of the action a from n to n'

Thus, a heuristic is consistent if, when going from neighboring nodes, the heuristic difference / step cost never overestimates the actual cost.

Lemmas

- If a heuristic is consistent, it is also admissible.
- If $h(n)$ is consistent, then the values of $f(n)$ on any path are non-decreasing.

3.3.1 Relaxed Problems

A relaxed problem is a problem that has fewer constraints on the actions than the original problem.

The cost of the optimal solution to a relaxed problem is an admissible heuristic for the original problem.

Example: Relaxed Problem as Admissible Heuristic

Consider a grid-based pathfinding problem where certain movements are restricted due to obstacles. A relaxed version of this problem might allow movement through obstacles by ignoring some constraints. The cost of the optimal solution to this relaxed problem, which is typically lower or equal to the original problem's optimal cost, serves as an admissible heuristic for the original problem.

This means that looking for relaxed problems is a good way to find admissible heuristics.

3.3.2 Dominance

A heuristic h_2 dominates h_1 if $h_2(n) \geq h_1(n)$ for all nodes n . (Given that h_1 and h_2 are admissible)

This means that a dominant heuristic is always closer to the optimal heuristic h^* , which results in less expansion and thus more efficient search.

3.3.3 Combining Heuristics

If h_1, h_2, \dots, h_m are admissible heuristics, then $h(n) = \max(h_1(n), h_2(n), \dots, h_m(n))$ is also admissible and dominates both h_1 and h_2 .

This is useful if we have multiple admissible non-dominated heuristics and want to combine them to get a better heuristic.

3.4 Informed Search

Definition

The informed search strategy has additional knowledge about "where" to look for solutions, usually in form of a heuristic function.


Some algorithms that utilize informed search are:

- Greedy Best-First Search
- A* Search
- Memory-Bounded Heuristic Search

3.4.1 Greedy Best-First Search

Greedy Best-First Search (GBFS) is an informed search algorithm that selects the next node to expand based on a heuristic that estimates the cost from the current node to the goal.

One example of a heuristic is the straight-line distance to the goal.



Example of GBFS with a straight-line distance heuristic.

Complexity: The time complexity of Greedy Best-First Search is $O(b^m)$, and its space complexity is also $O(b^m)$, where b is the branching factor and m is the maximum depth of the search tree. This is the same as DFS, however due to the heuristic it can be significantly faster.

GBFS is not guaranteed to be optimal or complete, especially if the heuristic function is not admissible or consistent. Its efficiency heavily depends on the quality of the heuristic used.


- **Complete?** No, can get stuck in loops.
- **Optimal?** No, does not guarantee an optimal solution as it depends on the heuristic, which may be flawed.

3.4.2 A* Search

Definition

A* Search is an informed search algorithm that combines the strengths of Uniform-Cost Search and Greedy Best-First Search. It searches for the least-cost path to the goal by considering both the cost to reach the current node and an estimated cost to reach the goal.

A* Search tries to avoid paths that are already expensive. It evaluates the complete path cost and the remaining cost to the goal. Its cost function is defined as $f(n) = g(n) + h(n)$, where $g(n)$ is the cost to reach the current node (root to n) and $h(n)$ (the heuristic) is the estimated cost to reach the goal (n to goal).



Example of A* Search with heuristic function.

Complexity: The time complexity of A* Search is $O(b^m)$, and its space complexity is also $O(b^m)$, where b is the branching factor and m is the maximum depth of the search tree.

Completeness and Optimality: A* Search is complete, and can be optimal, if the heuristic is admissible.

3.4.3 Alternatives to A*

1. Iterative-Deepening A* (IDA*)

- Like iterative deepening, it explores nodes level by level, but uses A* to evaluate the nodes.
- Cutoff information is the f-cost ($g + h$) instead of the depth.

2. Recursive Best-First Search (RBFS)

- Recursive algorithm that mimics best-first search with linear space
- Keeps track of f-value of best alternative path
- Path available from any ancestor of the current node and heuristic evaluations are updated with results of successors

3. (Simple) Memory-Bounded A* ((S)MA*)

- Drops the worst leaf node when memory is full
- Its value will be updated to its parent
- May be researched later

3.4.4 Graph Search

When traversing a problem, loops can occur. Failure to detect them can turn linear search into exponential search. To avoid this, we can use **graph search**. Hereby we only expand nodes that have not been visited yet.

For example: The graph search version of UCS is **Dijkstra's algorithm**.